

CS2040S

Data Structures and Algorithms

Hashing III

Hashing!

- Introduction to Hashing
- Collision Resolution: chaining
- Java hashing
- Collision resolution: open addressing
- Table (re)sizing

Midterm

Thurs. March 10 6:30pm

- Location: MPSH
- Room assignment on Coursemology
- Please double-check room

Bring to quiz:

- One double-sided sheet of paper with any notes you like.
- Pens/pencils.
- You may not use anything else. (No calculators, no phones, etc.)



Midterm

Covid Issues

- Please do test before taking the midterm.
- You will need to have the “green pass” on the Univus app to take midterm.
- Please do not come if you feel unwell.



Midterm

What happens if covid positive?

- Upload test to Univus.
- There will be a makeup next week.
- Get well soon!

What if I don't feel well?

- If covid positive, see above.
- If not covid positive, see doctor for MC → Makeup.
- But don't take midterm if unwell!



Hashing!

- Introduction to Hashing
- Collision Resolution: chaining
- Java hashing
- Collision resolution: open addressing
- Table (re)sizing

Abstract Data Types

Symbol Table

public interface SymbolTable


<code>void insert(Key k, Value v)</code>	<i>insert (k,v) into table</i>
<code>Value search(Key k)</code>	<i>get value paired with k</i>
<code>void delete(Key k)</code>	<i>remove key k (and value)</i>
<code>boolean contains(Key k)</code>	<i>is there a value for k?</i>
<code>int size()</code>	<i>number of (k,v) pairs</i>

Note: no successor / predecessor queries.

Hash Functions

Problem:

- Huge universe U of possible keys.
- Smaller number n of actual keys.
- How to map n keys to $m \approx n$ buckets?



Hash Functions

Collisions:

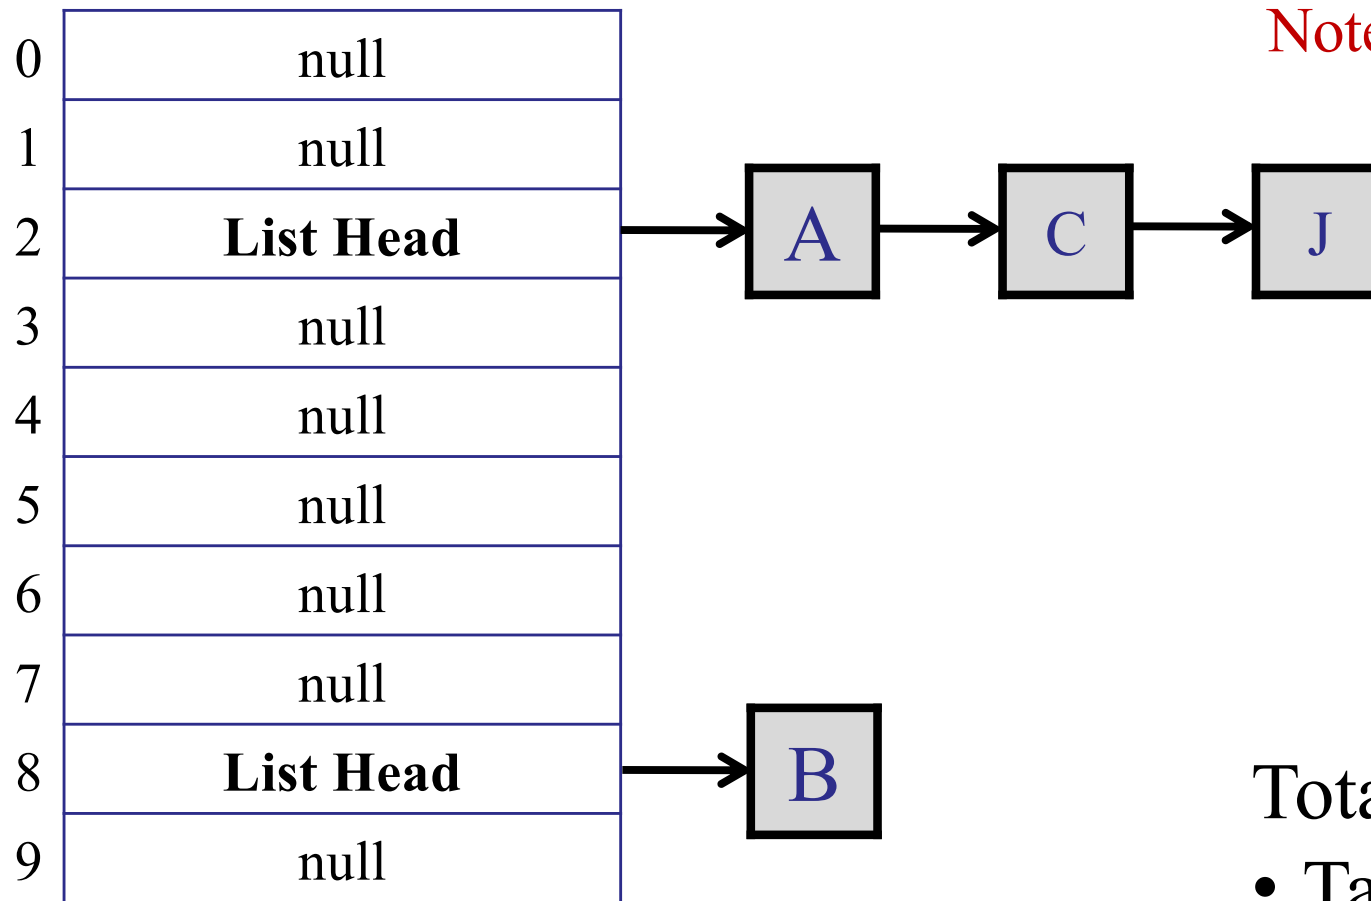
- We say that two distinct keys k_1 and k_2 **collide** if:

$$h(k_1) = h(k_2)$$

- Unavoidable!
 - The table size is smaller than the universe size.
 - The pigeonhole principle says:
 - There must exist two keys that map to the same bucket.
 - Some keys must collide!

Chaining

Each bucket contains a linked list of items.



Note: $h(A) == h(C) == h(J)$

Total space: $O(m + n)$

- Table size: m
- Linked list size: n

Let's be optimistic today.

The Simple Uniform Hashing Assumption

- Every key is equally likely to map to every bucket.
- Keys are mapped independently.

Assume hash function has this property, even if it may not!

Intuition:

- Each key is put in a random bucket.
- Then, as long as there are enough buckets, we won't get too many keys in any one bucket.

Hashing with Chaining

If hash function satisfies Simple Uniform Hashing Assumption

Searching:

- Expected search time = $1 + n/m = O(1)$
- Worst-case search time = $O(n)$

Inserting:

- Worst-case insertion time = $O(1)$

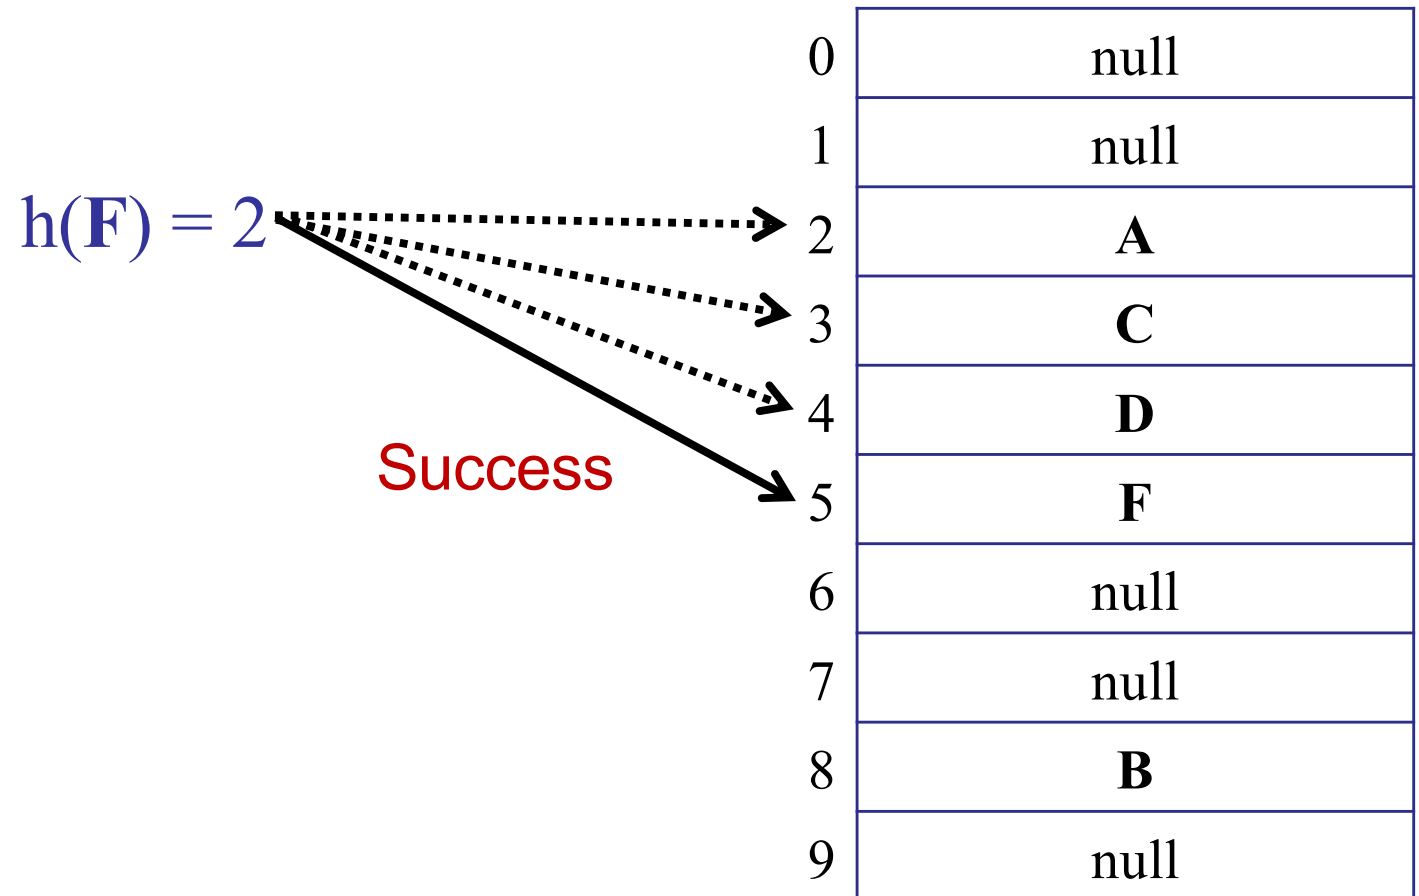
**** In this case, inserting allows duplicates...**

Preventing duplicates requires searching.

Open Addressing

On collision:

Probe a sequence of buckets until you find an empty one.



Hash Functions

Two properties of a good hash function:

1. $h(key, i)$ enumerates all possible buckets.
 - For every bucket j , there is some i such that:
$$h(key, i) = j$$
 - The hash function is permutation of $\{1..m\}$.
 - For linear probing: true!

Hash Functions

Two properties of a good hash function:

2. Uniform Hashing Assumption

Every key is equally likely to be mapped to every *permutation*, independent of every other key.

n! permutations for probe sequence: e.g.,


- 1 2 3 4
- 1 2 4 3
- 1 4 2 3
- 1 4 3 2
- ...

Performance of Open Addressing

- Chaining:
 - When $(m == n)$, we can still add new items to the hash table.
 - We can still search efficiently.
- Open addressing:
 - When $(m == n)$, the table is full.
 - We cannot insert any more items.
 - We cannot search efficiently.

Performance of Open Addressing

Define:

- Load $\alpha = n / m$  Average # items / bucket
- Assume $\alpha < 1$.

Claim:

For n items, in a table of size m , assuming *uniform hashing*, the expected cost of an operation is:


$$\leq \frac{1}{1 - \alpha}$$

Example: if ($\alpha=90\%$), then $E[\# \text{ probes}] = 10$

Disadvantages...

Open addressing:

- Performance degrades badly as $\alpha \rightarrow 1$.



Hashing: Recap

Problem: coping with large universe of keys

- Number of possible keys is very, very large.
- Direct Access Table takes too much space

Hash functions

- Use hash function to map keys to buckets.
- Sometimes, keys collide (inevitably!)

Resolve collisions

- Chaining \rightarrow SUHA $\rightarrow O(1 + \alpha)$ expected cost ops
- Open Addressing \rightarrow UHA $\rightarrow O(1 / 1 - \alpha)$ expected cost ops

Hashing!

- Introduction to Hashing
- Collision Resolution: chaining
- Java hashing
- Collision resolution: open addressing
- Table (re)sizing

Table Size

How large should the table be?

- Assume: Hashing with Chaining
- Assume: Simple Uniform Hashing
- Expected search time: $O(1 + n/m)$
- Optimal size: $m = \Theta(n)$
 - if $(m < 2n)$: too many collisions.
 - if $(m > 10n)$: too much wasted space.
- Problem: we don't know n in advance.

Table Size

Idea:

- Start with small (constant) table size.
- Grow (and shrink) table as necessary.

Example:

- Initially, $m = 10$.
- After inserting 6 items, table too small! Grow...
- After deleting $n-1$ items, table too big! Shrink...

Table Size

How to grow the table:

1. Choose new table size m .
2. Choose new hash function h .
 - Hash function depends on table size!
 - Remember: $h : U \rightarrow \{1..m\}$
3. For each item in the old hash table:
 - Compute new hash function.
 - Copy item to new bucket.

Table Size

Time complexity of growing the table:

— Assume:

- Let m_1 be the size of the old hash table.
- Let m_2 be the size of the new hash table.
- Let n be the number of elements in the hash table.

— Costs:

- Scanning old hash table: $O(m_1)$
- Inserting each element in new hash table: $O(1)$
- Total: $O(m_1 + n)$

Table Size

Time complexity of growing the table:

— Assume:

- Size $m_1 < n$.
- Size $m_2 > 2n$

— Costs:

- Total: $O(m_1 + n)$.
 $= O(n)$

Table Size

Time complexity of growing the table:

Wait! What is the cost of initializing the new table?

- Initializing a table of size X takes X time!

- Costs:

Total: $O(m_1 + m_2 + n)$

Table Size

Time complexity of growing the table:

— Assume:

- Let m_1 be the size of the old hash table.
- Let m_2 be the size of the new hash table.
- Let n be the number of elements in the hash table.

— Costs:

- Scanning old hash table: $O(m_1)$
- Creating new hash table: $O(m_2)$
- Inserting each element in new hash table: $O(1)$
- Total: $O(m_1 + m_2 + n)$

How fast to grow?

Idea 1: Increment table size by 1

- if ($n == m$): $m = m+1$

- Cost of resize:

- Size $m_1 = n$.
- Size $m_2 = n+1$.
- Total: $O(n)$

Initially: $m = 8$

Increase table size by 1 on resize.

What is the cost of inserting n items?

1. $O(n)$
2. $O(n \log n)$
- ✓ 3. $O(n^2)$
4. $O(n^3)$
5. None of the above.

ARCHIPELAGO

is open

How fast to grow?

Idea 1: Increment table size by 1

- When $(n == m)$: $m = m+1$
- Cost of each resize: $O(n)$

Table size	8	8	9	10	11	12	...	$n+1$
Number of items	0	7	8	9	10	11	...	n
Number of inserts		7	1	1	1	1	...	1
Cost		7	8	9	10	11		n

- Total cost: $(7 + 8 + 9 + 10 + 11 + \dots + n) = O(n^2)$

How fast to grow?

Idea 2: Double table size

- if $(n == m)$: $m = 2m$
- Cost of resize:
 - Size $m_1 = n$.
 - Size $m_2 = 2n$.
 - Total: $O(n)$

How fast to grow?

Idea 2: Double table size

- When $(n == m)$: $m = 2m$
- Cost of each resize: $O(n)$

Table size	8	8	16	16	16	16	16	16	16	16	32	32	32	...	2n
# of items	0	7	8	9	10	11	12	13	14	15	16	17	18	...	n
# of inserts		7	1	1	1	1	1	1	1	1	1	1	1	...	1
Cost		7	8	1	1	1	1	1	1	1	16	1	1		n

- Total cost: $(7 + 15 + 31 + \dots + n) = O(n)$

How fast to grow

Idea 2: Double table size

Cost of Resizing:

Table size	Total Resizing Cost
8	8
16	$(8 + 16)$
32	$(8 + 16 + 32)$
64	$(8 + 16 + 32 + 64)$
128	$(8 + 16 + 32 + 64 + 128)$
...	...
m	$<(1+2+4+8+\dots+m) \leq O(m)$

How fast to grow?

Idea 2: Double table size

- if ($n == m$): $m = 2m$
 - Cost of resize: $O(n)$
 - Cost of inserting n items + resizing: $O(n)$
- Most insertions: $O(1)$
- Some insertions: linear cost (expensive)
- Average cost: $O(1)$

Design question

Do you care that some insertions take a lot longer than others?

- Most insertions: $O(1)$
- Some insertions: linear cost (expensive)
- Total cost is good...
- ... but what if the slow operation is really, really important / time critical?
- What if YOUR online purchase is the one that triggers a two hour database rebuild?

How fast to grow?

Idea 3: Square table size

- When $(n == m)$: $m = m^2$

Table size	Total Resizing Cost
8	?
64	?
4,096	?
16,777,216	?
...	...
m	?

Assume: square table size

What is the cost of inserting n items?

1. $O(\log n)$
2. $O(\sqrt{n})$
3. $O(n)$
4. $O(n \log n)$
- ✓ 5. $O(n^2)$
6. $O(2^n)$
7. None of the above.

ARCHIPELAGO

is open

How fast to grow?

Idea 3: Square table size

- if $(n == m)$: $m = m^2$
- Cost of resize:
 - Size $m_1 = n$.
 - Size $m_2 = n^2$.
 - Total: $O(m_1 + m_2 + n)$
 $= O(n + n^2 + n)$
 $= O(n^2)$

How fast to grow?

Idea 3: Square table size

- When $(n == m)$: $m = m^2$

# Items	Total Resizing Cost
8	64
64	$(64 + 4,096)$
4,096	$(64 + 4,096 + \dots)$
...	...
n	$> n^2$
	$= O(n^2)$

How fast to grow?

Idea 3: Square table size

- When $(n == m)$: $m = m^2$

# Items	Resizing Cost	Insert Cost
8	64	8
64	$(64 + 4,096)$	64
4,096	$(64 + 4,096 + \dots)$	4,096
...
n	$> n^2$	n
	$< O(n^2)$	$O(n)$

How fast to grow?

Idea 3: Square table size

- if ($n == m$): $m = m^2$
- Cost of resize:
 - Total: $O(n^2)$
- Cost of inserts:
 - Total: $O(n)$

Why else is squaring the table size bad?

1. Resize takes too long to find items to copy.
- ✓ 2. Inefficient space usage.
3. Searching is more expensive in a big table.
4. Inserting is more expensive in big table.
5. Deleting is more expensive in a big table.

ARCHIPELAGO

is open

Deleting Elements

Basic procedure: (chained hash tables)

Delete(*key*)

1. Calculate hash of *key*.
2. Let *L* be the linked list in the specified bucket.
3. Search for item in linked list *L*.
4. Delete item from linked list *L*.

Cost:

- Total: $O(1 + n/m)$

Deleting Elements

What happens if too many items are deleted?

- Table is too big!
- Shrink the table...
- Try 1:
 - If $(n == m)$, then $m = 2m$.
 - If $(n < m/2)$ then $m = m/2$.

Deleting Elements

Rules for shrinking and growing:

– Try 1:

- If $(n == m)$, then $m = 2m$.
- If $(n < m/2)$ then $m = m/2$.

– Example problem:

- Start: $n=100, m=200$
- Delete: $n=99, m=200 \rightarrow$ shrink to $m=100$
- Insert: $n=100, m=100 \rightarrow$ grow to $m=200$
- Repeat...

Deleting Elements

Example execution:

- Start: $n=100$, $m=200$
- cost=100 • Delete: $n=99$, $m=200 \rightarrow$ shrink to $m=100$
- cost=100 • Insert: $n=100$, $m=100 \rightarrow$ grow to $m=200$
- cost=100 • Delete: $n=99$, $m=200 \rightarrow$ shrink to $m=100$
- cost=100 • Insert: $n=100$, $m=100 \rightarrow$ grow to $m=200$
- cost=100 • Delete: $n=99$, $m=200 \rightarrow$ shrink to $m=100$
- cost=100 • Insert: $n=100$, $m=100 \rightarrow$ grow to $m=200$
- Repeat...

Deleting Elements

Rules for shrinking and growing:

– Try 2:

- If $(n == m)$, then $m = 2m$.
- If $(n < m/4)$, then $m = m/2$.

– Claim:

- Every time you double a table of size m , at least $m/2$ new items were added.
- Every time you shrink a table of size m , at least $m/4$ items were deleted.

Amortized Analysis

Technique for analyzing “average” cost:

- Common in data structure analysis
- Smooths the cost when some operations are expensive and some operations are cheap:
 - E.g., some ops are $O(1)$, some are $O(n)$.
 - What matters is the *total* cost of all the ops.

Definition:


- Operation has amortized cost $T(n)$ if for every integer k , the cost of k operations is $\leq k T(n)$

Amortized Analysis

Definition:

- Operation has amortized cost $T(n)$ if for every integer k , the cost of k operations is $\leq k T(n)$

Example: amortized cost = 7



Amortized Analysis

“amortized” is NOT “average”

Definition:

- Operation has amortized cost $T(n)$ if for every integer k , the cost of k operations is $\leq k T(n)$

Example: amortized cost **NOT** 7

-
1. insert: 13 $13 > 7$
 2. insert: 5 $13+5 > 2*7 = 14$
 3. insert: 5 $13+5+5 > 3*7 = 21$
 4. insert: 5 $13+5+5+5 \leq 4*7 = 28$
 5. insert: 7 $5+5+5+13+7 \leq 5*7 = 35$

Amortized Analysis

Definition:

- Operation has amortized cost $T(n)$ if for every integer k , the cost of k operations is $\leq k T(n)$

Example: (Hash Tables)

- Inserting k elements into a hash table with resizing takes time $O(k)$.
- Conclusion:
The insert operation has amortized cost $O(1)$.

Amortized Analysis


Accounting Method

- Imagine a bank account **B**.
- Each operation adds money to the bank account.
- Every step of the algorithm spends money:
 - Immediate money: to perform the operation.
 - Deferred money: from the bank account.
- Total cost execution = total money
 - Average time / operation = money / num. ops

Amortized Analysis

Accounting Method Example (Hash Table)

- Each table has a bank account.
- Each time an element is added to the table, it adds $O(1)$ dollars to the bank account, uses $O(1)$ dollars to insert element.
- A table with k new elements since last resize has k dollars in bank.



0	null
1	null
2	(k_1, A)
3	null
4	null
5	null
6	null
7	null
8	(k_2, B)
9	null

Amortized Analysis

Accounting Method Example (Hash Table)

- Each table has a bank account.
- Each time an element is added to the table, it adds $O(1)$ dollars to the bank account.
- Claim:
 - Resizing a table of size m takes $O(m)$ time.
 - If you resize a table of size m , then:
 - at least $m/2$ new elements since last resize
 - bank account has $\Theta(m)$ dollars.

Amortized Analysis

Accounting Method Example (Hash Table)

- Each table has a bank account.
- Each time an element is added to the table, it adds $O(1)$ dollars to the bank account.
- Pay for resizing from the bank account!
- Strategy:
 - Analyze inserts ignoring cost of resizing.
 - Ensure that bank account always is big enough to pay for resizing.

Amortized Analysis

Total cost: Inserting k elements costs:

- Deferred dollars: $O(k)$ (to pay for resizing)
- Immediate dollars: $O(k)$ for inserting elements in table
- Total (Deferred + Immediate): $O(k)$

Amortized Analysis

Total cost: Inserting k elements costs:

- Deferred dollars: $O(k)$ (to pay for resizing)
- Immediate dollars: $O(k)$ for inserting elements in table
- Total (Deferred + Immediate): $O(k)$

Cost per operation:

- Deferred dollars: $O(1)$
- Immediate dollars: $O(1)$
- Total: $O(1)$ / per operation

Example: Binary Counter

Counter ADT:

- increment()
- read()

[illegible]

Example: Binary Counter

Counter ADT:

- increment()
- read()

increment()

[illegible]

Example: Binary Counter

Counter ADT:

- increment()
- read()

increment(), increment()

[illegible]

Example: Binary Counter

Counter ADT:

- increment()
- read()

increment(), increment(), increment()

[illegible]


What is the worst-case cost of incrementing a counter with max-value n ?

1. $O(1)$
- ✓ 2. $O(\log n)$
3. $O(n)$
4. $O(n^2)$
5. I have no idea.

Example: Binary Counter

Question: If we increment the counter to n , what is the amortized cost per operation?


- Easy answer: $O(\log n)$
- More careful analysis....



Example: Binary Counter

Observation:


During each increment, only one bit is changed
from: $0 \rightarrow 1$



Example: Binary Counter

Observation:

During each increment, many bits may be changed
from: $1 \rightarrow 0$




Example: Binary Counter

Observation:

Accounting method: each bit has a bank account.


Whenever you change it from $0 \rightarrow 1$, add one dollar.

Whenever you change it from $1 \rightarrow 0$, pay one dollar.



Example: Binary Counter

Counter ADT



A diagram showing a sequence of 10 black squares, each containing a white '0', representing a binary vector.


Example: Binary Counter

Counter ADT

increment()

0	0	0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---	---	---

0 0 0 0 0 0 0 0 0 0




0	0	0	0	0	0	0	0	0	1
---	---	---	---	---	---	---	---	---	---

0 0 0 0 0 0 0 0 0 1



Example: Binary Counter

Counter ADT

increment(), increment(), increment()



0 0 0 0 0 0 0 0 1 0




0 0 0 0 0 0 0 0 1 1



Example: Binary Counter

Counter ADT

increment()



0 1 1 1 1 1 1 1 1 1



1 0 0 0 0 0 0 0 0 0


Example: Binary Counter

Observation:

Amortized cost of increment: 2

- One operation to switch one $0 \rightarrow 1$
- One dollar (for bank account of switched bit).

(All switches from $1 \rightarrow 0$ paid for by bank account.)



Design question

Do you care that some insertions take a lot longer than others?

- Most insertions: $O(1)$
- Some insertions: linear cost (expensive)
- Total cost is good...
- ... but what if the slow operation is really, really important / time critical?
- What if YOUR online purchase is the one that triggers a two hour database rebuild?

Hash Table Resizing

Rules for shrinking and growing:

- If $(n == m)$, then $m = 2m$.
- If $(n < m/4)$, then $m = m/2$.

Key facts:

- Every time you double a table of size m , at least $m/2$ new items were added \rightarrow can pay for doubling.
- Every time you shrink a table of size m , at least $m/4$ items were deleted \rightarrow can pay for shrinking.

\rightarrow Operations cost $O(1)$ amortized, expected cost!

Hashing!

- Introduction to Hashing
- Collision Resolution: chaining
- Java hashing
- Collision resolution: open addressing
- Table (re)sizing